uild the room before you rite the memo. Grab the 4-rompt project room kit: ource inventory, duplicate g, missing-context list, rounded draft.

e first useful agent workflow is not generation. It is tting the work surface into shape.

NATE
MAY 22, 2026 · PAID

116 20 4 Share Transcript •••

hen AI produces a mediocre draft from a messy folder, e prompt is almost never the problem. The room is.

Nate's Notebook

Welcome to my podcast! In these audio reviews of my newsletters, I am to break down complex AI topics in a way that's approachable and relatable. I want you to walk away with the confidence to leverage AI more effectively at home and at work!

Listen on

Substack App

Apple Podcasts

e model has been handed a strategy doc, two slightly
ferent versions of the operating plan, a transcript
h two meetings in it, and a deck that no longer
tches reality. It is asked to write a memo. To do that,
as to do two jobs at once: figure out what the project
ually is, then produce the artifact. The first job is the
rd one. The second job is the one that shows up in the
ft.

sharper prompt won't fix this. You need to prepare the
om first.

ecently worked on a project where the real work did
t live in one place. A strategy doc, meeting
nscripts, a budget spreadsheet, trip-planning notes,
g-design drafts, old PDFs, follow-up emails, half-
ished notes. Some clearly current. Others superseded.
few useful only because they showed how the thinking
d changed.

e useful first prompt was much more boring than
rite the plan." It was something like: help me build
e room. Find the relevant materials. Preserve the
ginals. Make an inventory. I needed to know which
es were authoritative, which were duplicates, which
re old, which were missing. I asked it to summarize
ch source before synthesizing across them, and
licitly told it not to write the final deliverable yet.

ly after that did the writing prompt become simple.
e the current operating plan for the numbers, the
nscript for decision context, the older PDF only as

ckground, and flag unsupported claims rather than oothing them over. The room made those distinctions ible before the writing started.

is kind of workflow was not really available a year o. Agents could draft, summarize, and answer estions, but they were uneven at walking a folder tree, ening files in sequence, comparing dates across cuments, and inspecting metadata without losing the read. In the last few months that has changed. The rrent generation of agents is good at the small, boring, e-level operations the work actually requires. Which eans the bottleneck has moved. It is no longer “can the odel produce the artifact.” It is “is the source set in ape for the model to do anything useful with it.”

ere’s what’s inside:

- **What is an AI project room.** What a bounded workspace looks like for one serious job, and which tools to use for which source types.
- **Why AI fails with messy source files.** Why serious work fails when you skip the preparation step and jump straight to generation.
- **How to build an AI source inventory.** How to build the artifact that makes everything downstream inspectable.
- **Summaries, duplicates, and missing context.** The three preparation layers that prevent bad synthesis before it starts.

The writing prompt, once the room exists. What changes when you draft from a clean work surface instead of a raw file dump.

Grab the four prompts. A room-builder for file-system tools, an inventory-and-audit for uploaded docs, a grounded-draft prompt that cites every claim back to a source, and a refresh prompt for when new files arrive.

t's build the room.

Subscribers get all posts like these!

INK: Grab the Prompts

e hidden cost of skipping preparation is that you ver see which source the bad claim came from. The mo reads fine. The number is wrong. Three weeks er you find out the agent pulled from a draft someone perseded in February, and now you are unwinding a cision that was already made on bad data. These four ompts exist because that failure mode is invisible until s expensive. The room-builder and the inventory-and-dit force the agent to declare authority and surface nflicts before it drafts anything, so when something oks off you can trace it to a file instead of guessing. e grounded-draft prompt cites every claim back to a urce ID and flags anything the room does not support, hich means review happens against evidence rather an polish. The refresh prompt keeps the room honest

hen new files arrive, instead of letting an old working ef quietly outlive the project. Run them in sequence the messiest folder you have this week. The artifact u want by the end is not a draft. It is a room you can fend.

What is an AI project room

project room is a bounded workspace for one serious . It is smaller than a second brain, more specific than nowledge management system, and not the same oblem as building yourself a local AI computer. It is workspace for one project, set up so that an agent n do useful work inside it.

project room is the place where the agent gathers the terial for one piece of work and makes it legible fore anyone asks for a final answer.

r a consulting project, that might mean interview nscripts, client decks, data exports, prior proposals, d meeting notes. For a house purchase: inspection orts, disclosures, contractor estimates, mortgage cuments, and email threads. A Substack piece would ll source PDFs, transcripts, draft notes, product docs, reenshots, and prior related posts. A board memo eds the financial model, operating plan, old board ck, current KPI export, and the notes from the last ee review meetings.

e point is not to create a perfect archive. The point is create a usable work surface.

good project room separates originals from working tes and keeps file provenance intact. It creates an ventory, names the source hierarchy, identifies plicates, and produces short summaries of each urce. It flags missing context and tells you what is dy to use and what still needs human judgment.

ly after that should the agent write.

e tool choice matters. Use Claude Projects when you ed a bounded workspace with uploaded documents d reusable project instructions. Use ChatGPT ojects or ChatGPT file analysis for smaller source sets d spreadsheets. Use Cursor or Claude Code when the oject room includes code, technical docs, or a folder e that needs inspection. Use NotebookLM when the is research-heavy and source-bounded, especially en you want to work against many uploaded cuments rather than a live file system.

e workflow is overkill for a throwaway note. It is not erkill when the cost of being wrong exceeds the cost preparing the room: a board memo, hiring packet, estor update, diligence brief, regulatory response, al review, operating plan, or article that depends on urce coverage. Use the room when the final artifact ll travel farther than your ability to explain it.

ilding an agentic back-office pipeline is a different oblem entirely. If you are automating invoice ocessing, customer support triage, or any operational rkflow that runs on a schedule, you need a data ategy that lives upstream of any one project. The om is a unit of work, not a unit of infrastructure. Use for the deliverable you are preparing this week. Build e pipeline somewhere else.

Why AI fails with messy source les

ost people still treat AI like a file generator.

ey ask it to make a memo, build a spreadsheet, create eck, draft a proposal, or summarize a meeting. metimes that works. But in serious work, the hard rt is often not generating the first draft. The hard part getting the inputs into a state where generation is rth doing.

is is especially obvious in Office work. The dream rsion is "AI makes the spreadsheet" or "AI makes the werPoint." The real pain is different. The spreadsheet s messy exports, merged headers, hidden tabs, stale sumptions, formulas that look plausible but are wrong, d numbers that need to be tied back to a source. The ck has brand rules, executive norms, approved guage, old charts, new data, and claims that need dence. The problem is not whether AI can create a

e but whether it can handle the material around the e.

at same pattern applies to almost every knowledge oject.

hat you actually need isn't "write." You need a chain of aller operations: inspect, gather, normalize, reconcile, mmarize, verify, and only then produce. When you p those steps, the AI may still give you something ent. It may even look useful. But you have moved the iew burden downstream. Now you have to figure out ether the answer is grounded, whether the right urces were used, whether an old file contaminated the nclusion, and whether the agent missed the one cument that mattered.

u can change the entire interaction by shifting when e review happens. Review the source inventory front, not after the memo is done. Ask for authority nking upfront rather than discovering halfway rough that the agent used an obsolete deck. Request a nflict log and a missing-context list before the nthesis, not after. These moves prevent the polished emo from hiding where it came from.

ents shine when they're working on structure, the ntext, and judgment about what's in the room. That's ere they add real value.

ow to set up AI file organization

u saw this structure inside Prompt 1. Here's why it's aped this way.

e first job is to gather without destroying.

at means the agent should not delete files, overwrite ginals, or silently collapse duplicates without asking. useful file-system agent starts conservatively. It looks, pies when appropriate, records paths, and proposes tions before taking irreversible steps.

sually want a structure like this:

- 00_originals — untouched source files and preserved paths.
- 01_inbox — unsorted material that may or may not matter.
- 02_inventory — the table of sources, authority, relevance, and notes.
- 03_source_summaries — one short summary per important source.
- 04_working_brief — the synthesis layer before drafting.
- 05_outputs — drafts and finished artifacts.
- 99_review — duplicate logs, proposed moves, uncertainty lists, and human approval items.

is looks basic. It should. The folder structure gives e agent a safe place to operate. It also gives you a way supervise without reading every document first.

ow to build an AI source ventory

e most important artifact in this workflow is not the aft. It is the source inventory.

good source inventory changes the entire interaction. tells you what the agent thinks the project consists of. at gives you a chance to correct the working set fore the final artifact inherits the mistake.

minimum, the inventory should include file path, urce type, date, owner if known, relevance, authority el, current or superseded status, supported claims, itations, intended use in the final work, and notes for man review.

final signed agreement outranks a negotiation draft. e current spreadsheet beats a screenshot of last nth's numbers. A transcript may be more faithful to at was said, while a cleaned meeting note may be re useful for decisions and owners. The approved k represents the story even if the underlying data es somewhere else, and an old PDF may be useful ckground but not a source for current claims.

can help sort this out, but it should not be allowed to le the sorting. The inventory makes the judgment ible.

e trust habit is to review the inventory before thesis. Spot-check the files the agent marks as

thoritative. Open one source it calls current and one urce it calls superseded. Check whether dates, owners, d titles match what you know. If the agent cannot plain why one file outranks another, do not let it draft m the room yet.

r a writing project, this is especially important. A odel can synthesize across twenty sources in a way at feels smooth but erases the difference between nfirmed facts, reported claims, working assumptions, d your own interpretation. The inventory gives you a ce to say, "Use this file for the timeline, this one for e numbers, this transcript for the quote, and ignore at older draft except as background."

uch better than trying to debug the final prose.

ummaries before synthesis

e next step is source summaries.

is sounds like ordinary AI summarization, but the rpose is different. You are not asking for a general mmary because you are too busy to read. You are king the agent to prepare the source for later use.

useful project-room summary should answer five estions.

hat is this source?

hat does it contain that matters for the project?

hat claims, numbers, or decisions does it support?

hat are its limitations?

w should it be used in the final work?

at last question is important. A transcript gives you e and raw detail but not polished phrasing. A readsheet carries the current figures but not the ategic interpretation behind them. A slide deck presents the approved frame without the underlying dence. A screenshot proves that something was ible at a moment in time, but not that it remains true.

e agent should also pull out uncertainty. A garbled me in a transcript should be flagged, not guessed at. A e that appears to be a draft should be labeled as one. sagreements between documents should stay visible, t get smoothed over. Missing sources should be called t.

e discipline here is simple: do not let the agent come confident before the room is clean.

l duplicate file detection for rojects

ost people think duplicate detection is housekeeping. AI work, duplicates are a reasoning problem.

he agent sees three versions of a plan and does not ow which one is current, it may blend them. The same

nscript exported twice with different names gets erweighted. An old deck and a new deck with similar les become a source for wrong claims. A revised dget spreadsheet next to an earlier copy can produce eraged assumptions without any flag.

d synthesis starts here.

e fix is not to let the agent delete duplicates. The fix is make it report them.

good duplicate log should separate exact duplicates, ely duplicates, and version families. Exact duplicates pear identical. Likely duplicates share similar names, tes, or content. Version families represent the same ifact over time.

e agent should propose a current version, but it ould explain why. Newer modified date is not always ough. The current version might have “final” in the le, or it might be the version attached to a later email, the one whose contents match the latest meeting tes. Sometimes it is not possible to know without man review.

is is exactly the kind of work AI is good at. It can scan mes, contents, metadata, and surrounding context ickly. But the final authority decision often belongs to e human.

at division of labor is healthy. Let the agent find the ss. Do not let it silently resolve the mess when the

nsequences matter.

he missing-context list

e of the best signs that an AI agent is helping
perly is that it tells you what it does not have.

weak workflow produces an answer from whatever
ppens to be in front of it. A strong workflow produces
issing-context list before the answer.

r a serious project, the agent walks through the room
d notices three kinds of problems. First, what is
ssing: the decision no one documented, the number
th no source, the absent owner. Second, what is
biguous: which version is current, which draft is
tdated, where documents disagree. Third, what is
ngerous: unsupported claims, private sources that
ould not be used, inferences presented as facts.

at matters because the missing material is often more
portant than the available material. Your document
y say “as discussed” while the source of truth sits in
e discussion itself. A transcript may say “we agreed to
s,” but the actual decision may have changed later in
ck. The spreadsheet includes a number without the
sumptions. The deck includes a chart without the data
urce.

you ask for the final memo too early, these gaps
come hallucination traps. If you ask for the missing-
ntext list first, they become review items.

he writing prompt, once the oom exists

ce the project room exists, the writing prompt gets orter and better.

stead of saying:

Vrite a strategy memo from these files."

u say:

se the reviewed source inventory and working brief. eat the current operating plan as authoritative for mbers, the transcript as source material for decision ntext, and the older deck as background only. Draft e memo, cite claims back to source IDs, and flag ything not supported by the room."

e model is no longer guessing what the source set eans. You have already reviewed the room. You have eady decided which materials matter. The agent has eady summarized the evidence, identified conflicts, d listed gaps. Now the writing work can be grounded.

is is also better for review. When the draft says mething questionable, you can trace it back to the ventory or the source summary. Inferences should be eled. Unsupported claims should be flagged. When e source hierarchy changes, you can update the project om and regenerate from a cleaner base.

st-moving projects need maintenance. If new files ive daily, refresh the inventory before each serious afting pass. If the source hierarchy changes, update e working brief. A project that changes purpose needs ew room rather than a contaminated old one. The al is not a permanent archive but a clean enough work rface for the next decision.

u are not trying to make AI perfect. You are trying to ke its work inspectable.

ick one project this week

ock 45 minutes Tuesday afternoon. Pick the project ere you have the messiest folder and the most nsequential deliverable. Drop the relevant materials o Claude Projects, ChatGPT Projects, Cursor, or tebookLM depending on the source set. Use the ompts in the kit above. Do not let the agent synthesize til you have reviewed the inventory.

at single session will teach you more about how your n work is organized than another hour of prompt eaking. The artifact you want by the end is not a draft. s a source inventory, duplicate log, missing-context t, and working brief.

e learning ramp most people need is not “how do I ke the AI sound smarter?” but “how do I make the rk surface clean enough that AI can help?”

he shift from generation to reparation

e old AI question was whether the model could oduce the artifact.

n it write the memo? Make the spreadsheet? Create e deck? Summarize the transcript?

ose questions still matter, but they are no longer the st interesting ones. The more useful question is ether the agent can help prepare the conditions under ich good work happens.

n it find the right sources, tell which ones are current, eserve originals, and identify missing context instead inventing around it?

hink many people will first feel agents become useful ht there. Not as magical writers or autonomous ployees. Not as another demo where a model oduces a passable first draft from a clean prompt.

the assistant that walks into the messy room before e meeting starts, turns on the lights, labels the boxes, ds the missing folder, puts the current documents on e table, and tells you what still does not add up.

en you can write.

t before.

I make this Substack thanks to readers like you!

Learn about all my Substack tiers here

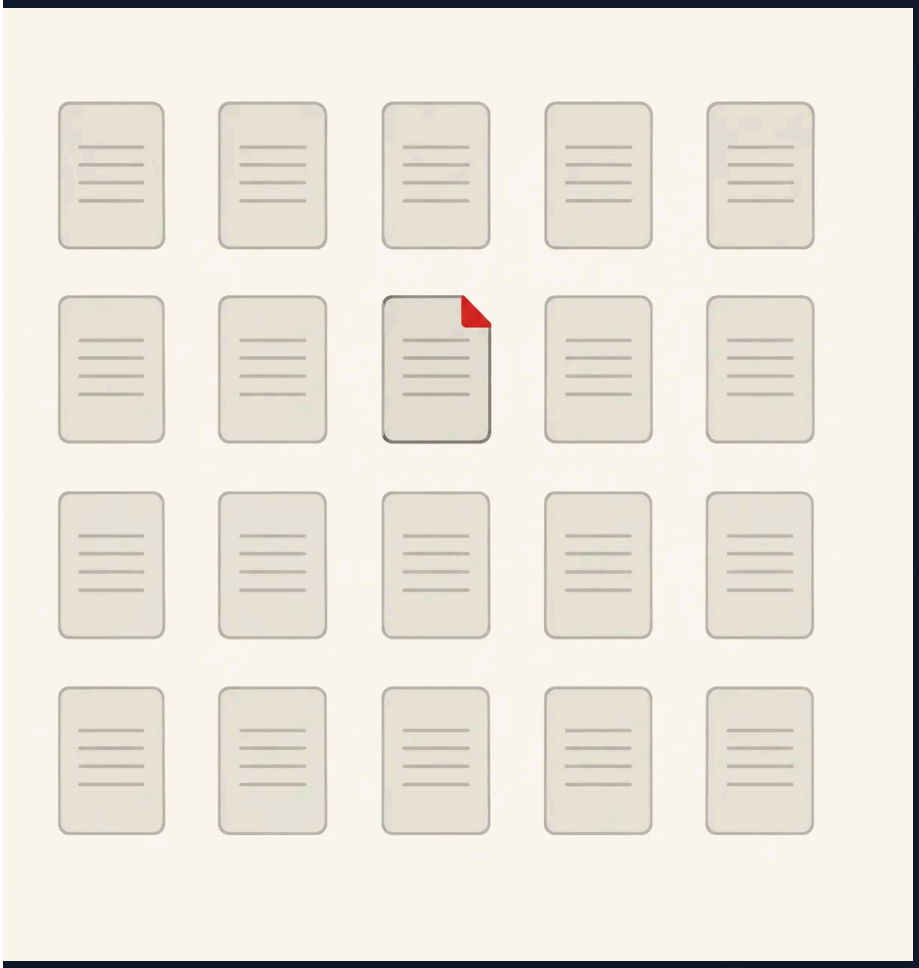

116 Likes · 4 Restacks

scussion about this video

omments Restacks

Write a comment...

Courtney Caldwell Eventual Eureka 2d

Man, this is one that really HITS HARD. What a common-sensical way to avoid the garbage in/garbage out (and all the invisible assumptions).

LIKE (7) REPLY SHARE

mikael thomas 2d *Edited*

That's very interesting, and it's exactly why I did not move further into production layers until I realized that you need a complete operational ontology, where the meaning and relationships between every element are declared within a structured context. From there, the system can grow systematically, becoming a magnet for discipline — stronger than any harness, though not necessarily replacing one.

If every angle of the system continuously maps back into this structured semantic framework, the model begins to perceive it as the most rewarding path, because it gains access to meaning, methods, skills, and intent as part of an integrated matrix of tools and relationships. At that point, it is no longer primarily about prompting, but about managing the model's attention through a declared symbolic structure.

This effectively channels the model's raw statistical inference into the domain context. That is where the alpha lives.

LIKE (3) REPLY SHARE

Kelly Heaton & ChatGPT-4o The Coherence Code 2d

This is why machine learning often fails to yield insights. The model cannot make meaning from raw data without an intelligent organizational structure.

LIKE (3) REPLY SHARE

Sr 2d

Loved this episode and it made me feel validated as I have been going in this direction with a similar approach I just blogged about https://stonebergdesign.com/blog/project-kickoff-skill I will be adopting the "data room" naming scheme

LIKE (2) REPLY (1) SHARE

Aaron Lawson Fork in the Road 1d

What are the function differences you see in your own project kickoff vs what Nate proposed?

LIKE REPLY (1) SHARE

Sr 1d

Nate talked about a file manifest or inventory with definitions as well as duplicate issues and I have not addressed those explicitly. They do get captured as a side benefit of the activity log though and that has been enough for me so far.

LIKE (1) REPLY SHARE

Thorsten 2d

This is timely insight as I'm attempting to build a research assistant with Claude inside of an Obsidian vault, using Zotero (and it's integration with Obsidian) to manage sources and citations. I'll attempt to start with Nate's proposed structure for each writing "project" and might at a later stage also try to work a Discourse Graph into the workflow.

LIKE (1) REPLY SHARE

Michael Jones Michael Jones 15h

SMH 🤦🏾 you did it again another tool for our toolbox 🧰 Class in session 🤫

LIKE REPLY SHARE

Aaron Lawson Fork in the Road 1d

As a bumbling vide coder moving at 90mph, I'm interested in hearing whether this schema makes sense for a new coding project repo, or if there are additional protocols that would make sense to swap out or add?

LIKE REPLY SHARE

The Geek in Review The Geek in Review 1d

I have been working on a project and applying (or so I thought) a lot of these principles in my workflow.

I had Claude 4.7 stop the project and review Nate's post and prompts. After one audit run, here's the first thing it said back "The audit produced findings I would have completely missed by proceeding to Phase E. Nate's framework just paid for itself. "

LIKE REPLY SHARE

mikael thomas 1d *Edited*

The useful first prompt was much more boring than "write the plan." It was something like: help me build the room. Find the relevant materials. Preserve the originals. Make an inventory. I needed to know which files were authoritative, which were duplicates, which were old, which were missing. I asked it to summarize each source before synthesizing across them, and explicitly told it not to write the final deliverable yet.

On this, i have completely generalised this process in three parts, the domain context that govern the agent behaviour and normalise all the static types elements, attached a redis memory for the current context in two layers, system and agents level, they can navigate my knowledge base and store what they need for the current job or consolidate my intent from reading transcripts of the past sessions, and the last that I just added that covers above is a store of object centric matrix. Like in this case observation matrix where I force the model to work on its one perception of the object, analyse dimensions, etc and I keep updating hit through until the project is done, then I store it as a type, then evaluate matrix or maintenance matrix etc.

This layer is portable cross agents and recoverable in time and extremely helpful to manage the model attention on one goal, especially when it's abstract.

LIKE REPLY SHARE

Robert LaSalle 2d

I just watched Nate's post today and the artifacts he suggests creating (like the file inventory) are gonna come in very handy when he discusses /goal in Claude Code...

https://code.claude.com/docs/en/goal

LIKE REPLY SHARE

Pamela Pierce 2d

I just spent a couple weeks creating the structure. Finally it's going to the AI for sorting.

LIKE REPLY SHARE

Alexander Verharen 2d

I've been building a solution that takes those original documents, working briefs and final docs to allow for truthful reporting. I built an entire digest-driven infrastructure for that. Every project has it's own pitfalls and my digest service provides preambles for prompts depending on your objective with these materials. It's very insightful to learn how I can get even more value out of this. I'd like to be able to do this with Notebook LM and tools like that

LIKE REPLY SHARE

Petko Mikhailov 2d

How can one make this work with Copilot Cowork in a Microsoft ecosystem (OneDrive, SharePoint, Teams, Loop, Power Automate)?

LIKE REPLY SHARE

Trtyptn 2d

I just built this as a 'Workstation' based on Jeff Yu's "Claude Cowork System". I will play with it this weekend on a side project I'm spinning up! @Nate - you left Cowork out of the equation, but I think Jeff's framework + your 'Clean Room' approach make a great combo!

Here's Jeff's video for reference
https://youtu.be/0_dSWLOHKng?si=t6ofgiAGvJaiPOxa

LIKE REPLY SHARE

Al Simon Al's Substack 2d

thanks nate, this is directly relevant to what i've been building - i had 5 parts of this already and your article gave me another 4.

just about to finish building it

LIKE REPLY SHARE

Saul Lowery Saul Lowery 2d

This is literally how I started building my Trust Fund Baby Architecture. All the way back to the first factory, you have to have a clean factory floor with everything in place, so the agent can do its job. Anything outside of that and you're setting it up for failure.

LIKE REPLY SHARE

John Halbert 2d *Edited*

funny, i turned to files as a means to make working with llms more reliable and effective early on, and here we are with files being first class in most harnesses in one way or another.

LIKE REPLY SHARE

Dina 2d

NOT BORING!

LIKE REPLY SHARE